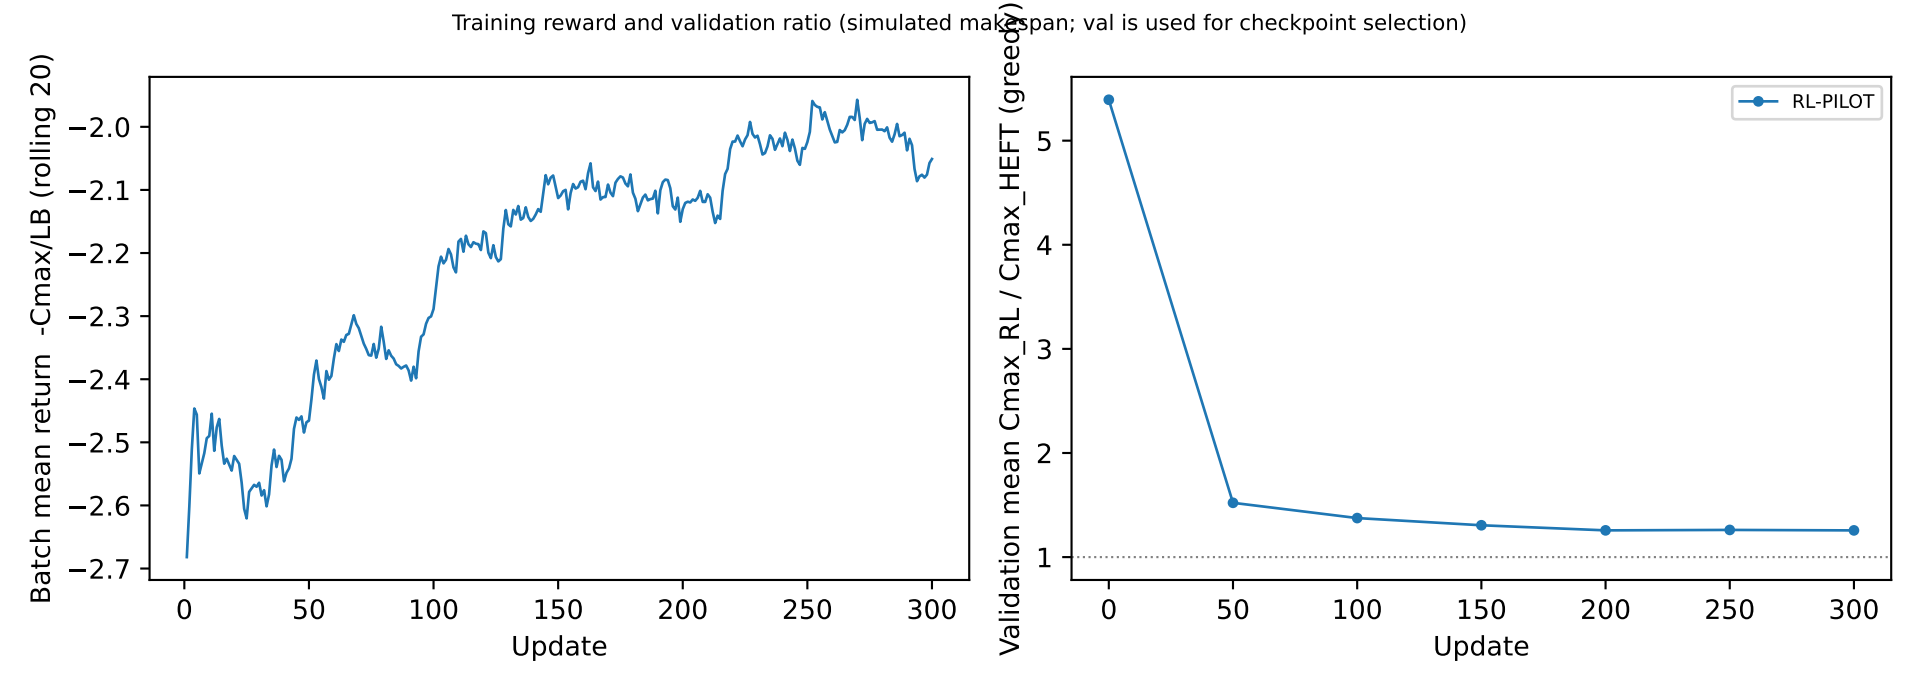
Training reward and validation ratio (simulated makespan; val is used for checkpoint selection)
Batch mean return  -Cmax/LB (rolling 20)
−2.0
−2.1
−2.2
−2.3
−2.4
−2.5
−2.6
−2.7
0
50
100
150
200
250
300
Update
Validation mean Cmax_RL / Cmax_HEFT (greedy)
RL-PILOT
5
4
3
2
1
0
50
100
150
200
250
300
Update